



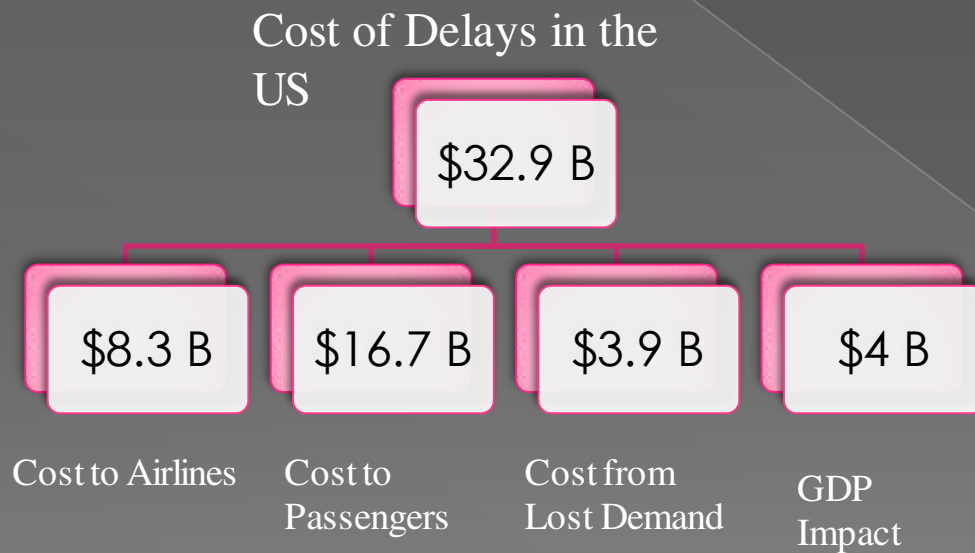
FLIGHT	DESTINATION	GATE #	STATUS
123	LOS ANGELES	A23	DELAYED
456	LONDON	C72	DELAYED
789	MADRID	B34	DELAYED
101	PARIS	A14	DELAYED
202	TOKYO	C89	DELAYED
303	HONG KONG	G12	DELAYED
404	MIAMI	C5	DELAYED
505	NEW YORK	D13	DELAYED
606	RIO DE JANEIRO	A4	DELAYED
707	SYDNEY	B22	DELAYED
808		A33	DELAYED

# USA AIRLINE DELAYS AND THEIR IMPACTS

Explained using Multiple Regression and  
Logistic Regression

-Satya M

# Why Airline Delays?



# How Different from Existing Models?

- ◉ I found 'N' number of projects in the internet trying to predict the Air Delay time based on various factors.
- ◉ However, I would focus on two important factors.
  - > AirTime
  - > WeatherDelay
- ◉ I'm using the above factors for consideration as the same hasn't been used before in any model and they are quite difficult to predict.
- ◉ They play an significant factor while looking at Air Line information.



# About the Dataset

- I used the dataset from Kaggle and it was populated using “The Flight Delay” project which is the most complete database of flight delays in the United States
- The dataset is separated into individual documents – one record for every year – and contains around 7 million records each.
- I’m considering the records only for the year 2008, as it holds the most recent data. Further, I’m looking at flight data with of Chicago and Washington DC.
- The source of data is given in the below link
- <https://www.kaggle.com/giovamata/airlinedelaycauses>

Carrier Delay



Carrier Delay



Weather Delay




Unique Carrier

AA

MQ


UA

NAS Delay




TaxiIn and TaxiOut vs Carrier


Unique Carrier



# What is Regression?



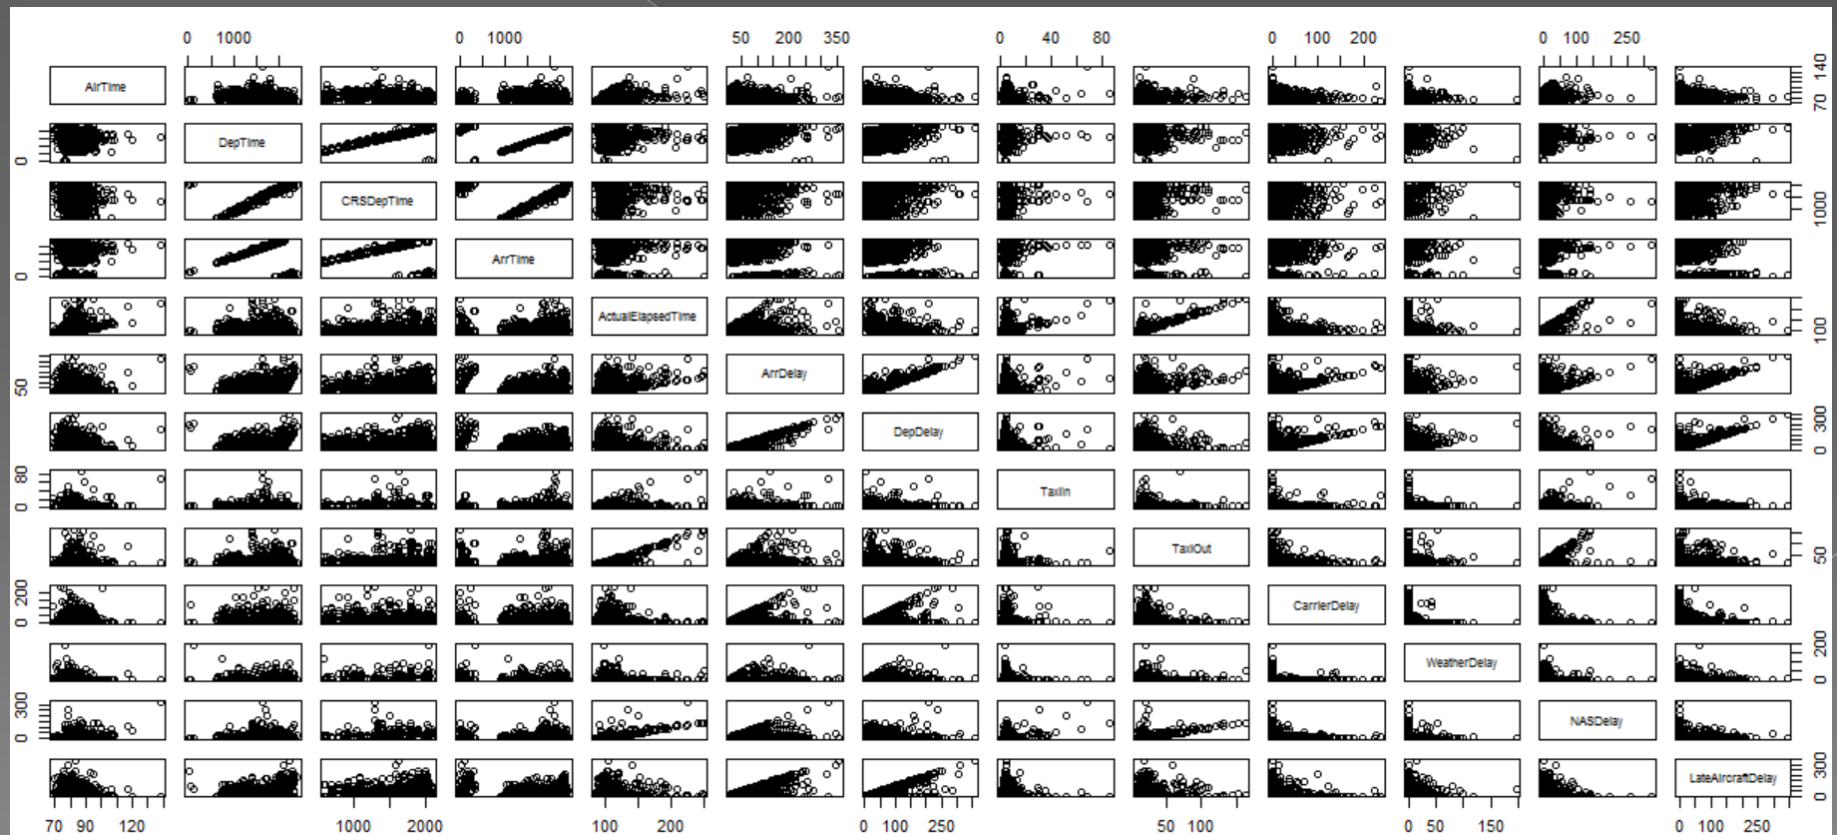
MY HOBBY: EXTRAPOLATING



- ❖ Regression Analysis is the art and science of fitting straight lines to patterns of data.
- ❖ Regression Analysis is widely used for predicting and forecasting.
- ❖ In a linear regression model, the variable of interest which is the dependent variable is predicted from a single (simple linear regression) or multiple group of independent variables (multiple linear regression).
- ❖ Whereas, for Logistic regression, the dependent variable is categorical (i.e. qualitative data)
- ❖ In this case, I included the following as the dependent variable.
  - ❖ For Multiple Regression – AirTime
  - ❖ For Logistic Regression – WeatherDelay  
(For Delay 0, else 1)

# Steps to Perform for Multiple Regression

- Draw a scatter plot and Observe whether you can fit a line to describe a pattern





# Steps to Perform for Multiple Regression continued...

- Build a multiple linear regression model ,  $y = f(x) = \beta_0 + \beta_1 X + e$
- ```
fit_model <- lm(model$AirTime ~  
as.factor(model$Month) +  
as.factor(model$DayofMonth) +  
as.factor(model$DayOfWeek) + model$DepTime +  
model$CRSDepTime + model$ArrTime +  
model$CRSArrTime + as.factor(model$UniqueCarrier)  
+ as.factor(model$FlightNum) +  
model$ActualElapsedTime + model$CRSElapsedTime  
+ model$ArrDelay + model$DepDelay + model$TaxiIn  
+ model$TaxiOut + model$CarrierDelay +  
model$WeatherDelay + model$NASDelay +  
model$LateAircraftDelay)
```



# Steps to Perform for Multiple Regression continued...

## Parameter Estimates

```
> # Dropping model$CarrierDelay
> fit_model <- lm(formula = model$AirTime ~
+               + model$ArrDelay + model$DepDelay +
+               + model$TaxiIn + model$TaxiOut)
> summary(fit_model)
```

Call:  
lm(formula = model\$AirTime ~ +model\$ArrDelay + model\$DepDelay +  
+model\$TaxiIn + model\$TaxiOut)

Residuals:

|  | Min     | 1Q      | Median  | 3Q     | Max    |
|--|---------|---------|---------|--------|--------|
|  | -7.1156 | -1.6715 | -0.3932 | 2.0813 | 8.4047 |

Coefficients:

|                 | Estimate  | Std. Error | t value | Pr(> t )   |
|-----------------|-----------|------------|---------|------------|
| (Intercept)     | 102.99754 | 0.38413    | 268.13  | <2e-16 *** |
| model\$ArrDelay | 0.89767   | 0.01442    | 62.26   | <2e-16 *** |
| model\$DepDelay | -0.90124  | 0.01424    | -63.28  | <2e-16 *** |
| model\$TaxiIn   | -0.82366  | 0.02253    | -36.56  | <2e-16 *** |
| model\$TaxiOut  | -0.87367  | 0.01522    | -57.40  | <2e-16 *** |

---  
Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 2.866 on 1015 degrees of freedom  
Multiple R-squared: 0.8025, Adjusted R-squared: 0.8018  
F-statistic: 1031 on 4 and 1015 DF, p-value: < 2.2e-16

# Steps to Perform for Multiple Regression continued...

## ● Goodness of Fit Test


- > A good model is determined by Goodness of fit.
- > There are 3 ways to determine the same.
- > These steps are to be performed sequentially.
  - 1. Overall Goodness using F-test
  - 2. Individual Parameter test
  - 3. Coefficient of determination ( $R^2$ )

# Steps to Perform for Multiple Regression continued...

## ● Residual Analysis

- > A residual plot is a scatter plot of the residuals against the predicted value.


**Residual = Actual value – Predicted value**



## Goals in Residual Analysis

- ❖ Validate the constant variance
- ❖ Validate the linearity relationship
- ❖ Validate normal distribution of residuals (QQ-Plot)
- ❖ Identify potential outliers (Cook's Distance)

```
#Residual Analysis
#Plot residuals vs predicted values
plot(fitted(fit_model), rstandard(fit_model), main="Predicted vs Residuals plot")
abline(a=0, b=0, col='red') #add zero line
```




Plot Residuals vs Predicted values

```
#Plot residuals vs each x-variables
fit_model

call:
lm(formula = model$AirTime ~ +model$ArrDelay + model$DepDelay +
+model$TaxiIn + model$TaxiOut)


Coefficients:
(Intercept)    model$ArrDelay    model$DepDelay    model$TaxiIn    model$TaxiOut
 102.9975         0.8977        -0.9012        -0.8237        -0.8737

#Plot residuals vs ArrDelay
plot(model$ArrDelay, rstandard(fit_model), main="ArrDelay vs residuals plot")
abline(a=0, b=0, col='red') #add zero line
```




Plot Residuals vs ArrDelay

```
#Plot residuals vs DepDelay
plot(model$DepDelay, rstandard(fit_model), main="DepDelay vs residuals plot")
abline(a=0, b=0, col='red') #add zero line
```




Plot Residuals vs DepDelay

```
#Plot residuals vs TaxiIn
plot(model$TaxiIn, rstandard(fit_model), main="TaxiIn vs residuals plot")
abline(a=0, b=0, col='red') #add zero line
```



Plot Residuals vs TaxiIn


```
#Plot residuals vs TaxiOut
plot(model$TaxiOut, rstandard(fit_model), main="TaxiOut vs residuals plot")
abline(a=0, b=0, col='red') #add zero line
```



Plot Residuals vs TaxiOut

# QQ Plot and Outliers

```
> # QQ-plot
> qq_plot <- (rstandard(fit_model))
> qqnorm(qq_plot)
> qqline(rstandard(fit_model), col='blue')
```



```
> #Influential Points
> #Cook's Distance
> # n=1020, 4/n = 4/1020 = 0.003
> cook <- influence.measures(fit_model)
> summary(cook)
Potentially influential observations of
lm(formula = model$AirTime ~ model$ArrDelay + model$DepDelay + model$TaxiIn + model$Taxiout) :
```

|     | dfb.1_ | dfb.m\$AD | dfb.m\$DD | dfb.m\$TI | dfb.m\$TO | dffit   | cov.r  | cook.d | hat    |
|-----|--------|-----------|-----------|-----------|-----------|---------|--------|--------|--------|
| 53  | 0.01   | 0.00      | 0.00      | 0.00      | 0.00      | -0.03   | 1.04_* | 0.00   | 0.03_* |
| 63  | 0.04   | 0.01      | -0.01     | 0.01      | -0.06     | -0.15   | 1.03_* | 0.00   | 0.03_* |
| 66  | 0.00   | 0.00      | 0.00      | 0.00      | 0.00      | 0.02    | 1.02_* | 0.00   | 0.01   |
| 72  | 0.00   | 0.00      | 0.00      | 0.00      | 0.00      | 0.01    | 1.01_* | 0.00   | 0.01   |
| 103 | 0.05   | -0.03     | 0.02      | -0.22     | 0.01      | -0.36_* | 1.00   | 0.03   | 0.02_* |
| 115 | 0.02   | -0.02     | 0.00      | 0.04      | 0.03      | -0.21   | 1.01   | 0.01   | 0.02_* |
| 121 | 0.01   | -0.03     | 0.03      | 0.04      | -0.04     | -0.20   | 1.04_* | 0.01   | 0.03_* |

# Regression Diagnostics

## ● $R^2$ / Adj - $R^2$ Coefficient of Determination

- > This is a measure of goodness of fit for a linear regression model.
- > Coefficient of determination is the calculation of the variation in the dependent variable to the variation in the independent variable. 0 means no linear relationship and 1 means perfect model. Usually  $R^2$  value  $> 80\%$  is considered a good model.
- > From the output, it is evident that 80.18% variation in Y is explained by X.

```
Residual standard error: 2.866 on 1015 degrees of freedom  
Multiple R-squared: 0.8025, Adjusted R-squared: 0.8018
```

- > The only difference between  $R^2$  and Adjusted  $R^2$  is that adjusted  $R^2$  increases only when the new term added improves the model. Hence, it is more reliable than  $R^2$

# Taboo on $R^2$

- High  $R^2$  means a better model ?
- Low  $R^2$  means a bad model ?
- I found the above taboos a myth, since post execution of my model, I'm able to find Adj-  $R^2 = 1$  at a point. I had to recheck my X variables.
- Since they were highly collinear, I got 1 as the Adj-  $R^2$  value, which is indeed not correct.





# Multicollinearity



- ❖ Multicollinearity is an undesirable situation where the correlations among the independent variables are strong.
- ❖ When two X variables are highly collinear, be it negative or positive, they essentially convey the same information and when this happens, I found the regression results to be paradoxical.
- ❖ Situations that indicate Multicollinearity:
  - High F-test, but none of the X variables are significant.
  - Appearance of NA in the parameter estimates.

# Problems due to Multicollinearity

- Multicollinearity misleadingly inflates the standard errors of coefficients.
- Thus, it makes some variables statistically insignificant while they should be otherwise significant.
- It is like two or more people are singing loudly at the same time. One cannot discern which is which. They offset each other.



# How to detect Multicollinearity?

- ❖ Variation Inflation Factors (VIF)  $\geq 5$  indicates Multicollinearity.
- ❖ If there are two or more variables that has VIF greater than or around 5, one of the variables must be removed first.
- ❖ To determine the best one to remove, remove each one individually.
- ❖ Select the regression equation with highest  $R^2$ .



# Steps to Perform for Multiple Regression continued...

- Evaluations and Predictions
  - > Hold Out Evaluation (Large Data)
  - > N-Fold Evaluation (Small Data)
- I performed Hold Out Evaluation.



# Model Selection

- Search Algorithms

- > Best Subset Regression
- > Backward Elimination
- > Stepwise Regression/Forward Selection

- Model Selection Methods

- > Adj- $R^2$
- > Mallows's  $C_p$  Statistics
- > AIC and BIC criterion
- > PRESS Statistics

```

> #Selecting the best model based on Adj-R2
> summary(fit_model)

Call:
lm(formula = model$AirTime ~ ++model$ArrDelay + model$DepDelay +
    ++model$TaxiIn + model$TaxiOut)

Residuals:
    Min       1Q   Median       3Q      Max
-7.1156 -1.6715 -0.3932  2.0813  8.4047

Coefficients:
              Estimate Std. Error t value Pr(>|t|)
(Intercept)   102.99754    0.38413   268.13  <2e-16 ***
model$ArrDelay    0.89767    0.01442    62.26  <2e-16 ***
model$DepDelay   -0.90124    0.01424   -63.28  <2e-16 ***
model$TaxiIn     -0.82366    0.02253   -36.56  <2e-16 ***
model$TaxiOut    -0.87367    0.01522   -57.40  <2e-16 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 2.866 on 1015 degrees of freedom
Multiple R-squared:  0.8025,    Adjusted R-squared:  0.8018
F-statistic: 1031 on 4 and 1015 DF,  p-value: < 2.2e-16

> #As Backward elimination has high Adj-R2(80.18%), so we take this model into consideration
> summary(fit_model2) #final reduced model by forward selection

Call:
lm(formula = model$AirTime ~ model$ActualElapsedTime + model$DepDelay +
    model$TaxiOut + model$ArrDelay)

Residuals:
    Min       1Q   Median       3Q      Max
-42.783  -1.536  -0.052   1.799  17.442

Coefficients:
              Estimate Std. Error t value Pr(>|t|)
(Intercept)   21.66587    4.35185    4.979 7.52e-07 ***
model$ActualElapsedTime  0.67723    0.04124   16.422 < 2e-16 ***
model$DepDelay    0.12173    0.04268    2.852  0.00443 **
model$TaxiOut     -0.54940    0.01590  -34.550 < 2e-16 ***
model$ArrDelay    -0.13474    0.04261   -3.162  0.00161 **
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 3.878 on 1015 degrees of freedom
Multiple R-squared:  0.6386,    Adjusted R-squared:  0.6372
F-statistic: 448.3 on 4 and 1015 DF,  p-value: < 2.2e-16

```



```
> #Backward - Stepwise
> step(fit_model2, direction="backward", trace=TRUE)
Start: AIC=2769.58
model$AirTime ~ model$ActualElapsedTime + model$DepDelay + model$TaxiOut +
  model$ArrDelay
```

|                            | Df | Sum of Sq | RSS   | AIC    |
|----------------------------|----|-----------|-------|--------|
| <none>                     |    |           | 15261 | 2769.6 |
| - model\$DepDelay          | 1  | 122.3     | 15383 | 2775.7 |
| - model\$ArrDelay          | 1  | 150.3     | 15411 | 2777.6 |
| - model\$ActualElapsedTime | 1  | 4054.5    | 19315 | 3007.9 |
| - model\$TaxiOut           | 1  | 17947.0   | 33208 | 3560.6 |

```
Call:
lm(formula = model$AirTime ~ model$ActualElapsedTime + model$DepDelay +
  model$TaxiOut + model$ArrDelay)
```

```
Coefficients:
      (Intercept)  model$ActualElapsedTime  model$DepDelay  model$TaxiOut  model$ArrDelay
      21.6659          0.6772          0.1217        -0.5494        -0.1347
```

```
> > step(fit_model1, scope=list(upper=fit_model2, lower=~1), direction="forward", trace=TRUE)
Start: AIC=3619.02
model$AirTime ~ model$ActualElapsedTime
```

|                   | Df | Sum of Sq | RSS   | AIC    |
|-------------------|----|-----------|-------|--------|
| + model\$TaxiOut  | 1  | 19493.0   | 15809 | 2801.6 |
| + model\$ArrDelay | 1  | 982.0     | 34320 | 3592.2 |
| + model\$DepDelay | 1  | 857.7     | 34444 | 3595.9 |
| <none>            |    |           | 35302 | 3619.0 |

```
Step: AIC=2801.6
model$AirTime ~ model$ActualElapsedTime + model$TaxiOut
```

|                   | Df | Sum of Sq | RSS   | AIC    |
|-------------------|----|-----------|-------|--------|
| + model\$ArrDelay | 1  | 426.14    | 15383 | 2775.7 |
| + model\$DepDelay | 1  | 398.13    | 15411 | 2777.6 |
| <none>            |    |           | 15809 | 2801.6 |

```
Step: AIC=2775.73
model$AirTime ~ model$ActualElapsedTime + model$TaxiOut + model$ArrDelay
```

|                   | Df | Sum of Sq | RSS   | AIC    |
|-------------------|----|-----------|-------|--------|
| + model\$DepDelay | 1  | 122.31    | 15261 | 2769.6 |
| <none>            |    |           | 15383 | 2775.7 |

```
Step: AIC=2769.58
model$AirTime ~ model$ActualElapsedTime + model$TaxiOut + model$ArrDelay +
  model$DepDelay
```

```
Call:
lm(formula = model$AirTime ~ model$ActualElapsedTime + model$TaxiOut +
  model$ArrDelay + model$DepDelay)
```

```
Coefficients:
```



# Model Evaluation

## ● Hold Out Evaluation

```
> y1=predict.glm(fit_model,test.data)
Warning message:
'newdata' had 204 rows but variables found have 1020 rows
>
> y2=predict.glm(fit_model2,test.data)
Warning message:
'newdata' had 204 rows but variables found have 1020 rows
> y=test.data[,8]
> rmse1 <- sqrt((y-y1)%*%(y-y1))/nrow(test.data)
> rmse2 <- sqrt((y-y2)%*%(y-y2))/nrow(test.data)
> rmse1
      [,1]
[1,] 1.34746
> rmse2
      [,1]
[1,] 1.28515
> |
```

# Measuring Predictive Performance

Root Mean Square Error :

*Best model minimizes RMSE*

$$RMSE = \sqrt{\frac{\sum_{i=1}^m (y_i - \hat{y}_i)^2}{m}}$$

Mean Absolute Error

*Best model minimizes MAE*

$$MAE = \frac{\sum_{i=1}^m |y_i - \hat{y}_i|}{m}$$

# Logistic Regression

- Predictor – Categorical or Numeric
- Response – Categorical
- Relationship between response(binary) and predictor(s)

Model for probability  $p = \Pr(Y=1)$ :

$$\log\left(\frac{p}{1-p}\right) = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \beta_3 x_3 + \beta_4 x_4 + e$$

# Steps to perform Logistic Regression

- Based on AIC or BIC value
- In this case, AIC value is used

```
> #Try with backward selection model
> step(full,direction = "backward",trace = F)

Call: glm(formula = train_data$Wdelay ~ train_data$Month + train_data$DayofMonth +
  train_data$UniqueCarrier + train_data$FlightNum + train_data$ActualElapsedTime +
  train_data$CRSElapsedTime + train_data$ArrDelay + train_data$CarrierDelay +
  train_data$NASDelay, family = "binomial", data = train_data,
  control = list(maxit = 50))

Coefficients:
              (Intercept)              train_data$Month              train_data$DayofMonth
              3.5873941                -0.4126614                -0.0202636
  train_data$UniqueCarrierMQ  train_data$UniqueCarrierUA  train_data$FlightNum
              -2.5555662                -0.6128026                0.0005114
  train_data$ActualElapsedTime  train_data$CRSElapsedTime  train_data$ArrDelay
              0.0797740                -0.1159837                0.0100691
  train_data$CarrierDelay      train_data$NASDelay
              -0.0361235                -0.0844027

Degrees of Freedom: 815 Total (i.e. Null); 805 Residual
Null Deviance: 637.8
Residual Deviance: 498.8      AIC: 520.8
```

# Steps to perform Logistic Regression contd.

```
> full1 <- glm(train_data$Wdelay ~ train_data$Month+train_data$ArrDelay+train_data$DepDelay+train_data$CarrierDelay+train_data$NASDelay,data = train_data, family = "binomial", control = list(maxit = 50))
> summary(full1)
```

Call:

```
glm(formula = train_data$Wdelay ~ train_data$Month + train_data$ArrDelay +
    train_data$DepDelay + train_data$CarrierDelay + train_data$NASDelay,
    family = "binomial", data = train_data, control = list(maxit = 50))
```

Deviance Residuals:

| Min     | 1Q      | Median  | 3Q      | Max    |
|---------|---------|---------|---------|--------|
| -2.0153 | -0.5361 | -0.3468 | -0.1928 | 3.2395 |

Coefficients:

|                          | Estimate  | Std. Error | z value | Pr(> z )     |
|--------------------------|-----------|------------|---------|--------------|
| (Intercept)              | -0.430183 | 0.291422   | -1.476  | 0.139904     |
| train_data\$Month        | -0.433393 | 0.071232   | -6.084  | 1.17e-09 *** |
| train_data\$ArrDelay     | 0.081100  | 0.021615   | 3.752   | 0.000175 *** |
| train_data\$DepDelay     | -0.071429 | 0.021580   | -3.310  | 0.000933 *** |
| train_data\$CarrierDelay | -0.036328 | 0.008238   | -4.410  | 1.03e-05 *** |
| train_data\$NASDelay     | -0.073820 | 0.022929   | -3.219  | 0.001284 **  |

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 637.84 on 815 degrees of freedom  
Residual deviance: 521.19 on 810 degrees of freedom  
AIC: 533.19

Number of Fisher Scoring iterations: 6

# Steps to perform Logistic Regression contd.

## ● McFadden's R-square

```
> mcFmodel <- glm(formula = train_data$Wdelay ~ train_data$Month + train_data$DayofMonth + train_data$UniqueCarrier + train_data$FlightNum + train_data$ActualElapsedTime + train_data$CRSElapsedTime + train_data$ArrDelay + train_data$CarrierDelay + train_data$NASDelay, family = "binomial", data = train_data, control = list(maxit = 50))
> mcFnmodel <- glm(train_data$Wdelay ~ 1, family = "binomial")
> 1-logLik(mcFmodel)/logLik(mcFnmodel)
'log Lik.' 0.2179834 (df=11)
> |
```

## ● 95% CI for the coefficients

```
> #Confidence Intervals For Model Coefficients
> confint(mcFmodel)
Waiting for profiling to be done...
```

|                               | 2.5 %         | 97.5 %       |
|-------------------------------|---------------|--------------|
| (Intercept)                   | -7.0824671151 | 14.272594321 |
| train_data\$Month             | -0.5691316353 | -0.263144833 |
| train_data\$DayofMonth        | -0.0452549713 | 0.004323568  |
| train_data\$UniqueCarrierMQ   | -5.0361771094 | -0.315940003 |
| train_data\$UniqueCarrierUA   | -1.2386408595 | 0.023931861  |
| train_data\$FlightNum         | -0.0001104443 | 0.001136609  |
| train_data\$ActualElapsedTime | 0.0366785209  | 0.126898499  |
| train_data\$CRSElapsedTime    | -0.2250645625 | -0.008808043 |
| train_data\$ArrDelay          | 0.0057002817  | 0.014555780  |
| train_data\$CarrierDelay      | -0.0539527266 | -0.021690803 |
| train_data\$NASDelay          | -0.1349492785 | -0.038537181 |

# Steps to perform Logistic Regression contd.

- Computing  $\exp(\text{coefficients})$  to analyze change in odds for changes in X

```
> exp(coef(mcfmodel))
```

|                               |                             |                        |
|-------------------------------|-----------------------------|------------------------|
| (Intercept)                   | train_data\$Month           | train_data\$DayofMonth |
| 36.13977707                   | 0.66188636                  | 0.97994035             |
| train_data\$UniqueCarrierMQ   | train_data\$UniqueCarrierUA | train_data\$FlightNum  |
| 0.07764825                    | 0.54183018                  | 1.00051151             |
| train_data\$ActualElapsedTime | train_data\$CRSElapsedTime  | train_data\$ArrDelay   |
| 1.08304225                    | 0.89048976                  | 1.01011998             |
| train_data\$CarrierDelay      | train_data\$NASDelay        |                        |
| 0.96452116                    | 0.91906105                  |                        |



# Steps to perform Logistic Regression contd.

## ● Change in Odds

```
> exp(confint(mcfModel))  
Waiting for profiling to be done...
```

|                               |  | 2.5 %       | 97.5 %       |
|-------------------------------|--|-------------|--------------|
| (Intercept)                   |  | 0.000839699 | 1.579461e+06 |
| train_data\$Month             |  | 0.566016734 | 7.686306e-01 |
| train_data\$DayofMonth        |  | 0.955753761 | 1.004333e+00 |
| train_data\$UniqueCarrierMQ   |  | 0.006498544 | 7.291032e-01 |
| train_data\$UniqueCarrierUA   |  | 0.289777799 | 1.024221e+00 |
| train_data\$FlightNum         |  | 0.999889562 | 1.001137e+00 |
| train_data\$ActualElapsedTime |  | 1.037359478 | 1.135302e+00 |
| train_data\$CRSElapsedTime    |  | 0.798464666 | 9.912306e-01 |
| train_data\$ArrDelay          |  | 1.005716559 | 1.014662e+00 |
| train_data\$CarrierDelay      |  | 0.947476896 | 9.785428e-01 |
| train_data\$NASDelay          |  | 0.873760229 | 9.621959e-01 |

# Steps to perform Logistic Regression contd.

## ● VIF

```
> vif(mcfmodel)
```

|       | GVIF      | Df | GVIF <sup>1/(2*Df)</sup> |
|-------|-----------|----|--------------------------|
| mon   | 1.179723  | 1  | 1.086150                 |
| dom   | 1.020857  | 1  | 1.010375                 |
| uq    | 3.713334  | 2  | 1.388165                 |
| fn    | 3.336038  | 1  | 1.826482                 |
| aet   | 17.387084 | 1  | 4.169782                 |
| crset | 1.330932  | 1  | 1.153660                 |
| ad    | 1.172485  | 1  | 1.082813                 |
| cd    | 1.057806  | 1  | 1.028497                 |
| nasd  | 16.854155 | 1  | 4.105381                 |

# Steps to perform Logistic Regression contd.

- Sensing how strong the predictor is

```
>
> set.seed(1021)
> n <- 1020
> x <- 1*(runif(n)<0.5)
> pr <- (x==1)*0.9+(x==0)*0.1
> y <- 1*(runif(n) < pr)
> mod <- glm(y~x, family="binomial")
> nullmod <- glm(y~1, family="binomial")
> 1-logLik(mod)/logLik(nullmod)
'log Lik.' 0.4928337 (df=2)
> |
```

```
> set.seed(1021)
> n <- 1020
> x <- 1*(runif(n)<0.5)
> pr <- (x==1)*0.99+(x==0)*0.01
> y <- 1*(runif(n) < pr)
> mod <- glm(y~x, family="binomial")
> nullmod <- glm(y~1, family="binomial")
> 1-logLik(mod)/logLik(nullmod)
'log Lik.' 0.8893839 (df=2)
> |
```

```
> set.seed(1021)
> n <- 1020
> x <- 1*(runif(n)<0.5)
> x <- 1*(runif(n)>0.5)
> pr <- (x==1)*0.9+(x==0)*0.1
> y <- 1*(runif(n) < pr)
> mod <- glm(y~x, family="binomial")
> nullmod <- glm(y~1, family="binomial")
> 1-logLik(mod)/logLik(nullmod)
'log Lik.' 0.4765672 (df=2)
> |
```

# Steps to perform Logistic Regression contd.

## ○ Predicting values based on equation

```
> pred_dataglm <- data.frame(mon=6,dm=30,uq="AA",fn=1448,aet=101,crset=105,ad=39,cd=10,nasd=0)
> predict(mcfmodel, pred_dataglm,se.fit = TRUE,interval=c("none","confidence","prediction"), level=0.95,type="response")
$fit
      1
0.05490634

$se.fit
      1
0.02068382

$residual.scale
[1] 1

> predict(mcfmodel, pred_dataglm, type="response")
      1
0.05490634
> pred_dataglm <- data.frame(mon=6,dm=5,uq="AA",fn=580,aet=98,crset=105,ad=120,cd=0,nasd=0)
> predict(mcfmodel, pred_dataglm, type="response")
      1
0.1364101
> pred_dataglm <- data.frame(mon=1,dm=3,uq="UA",fn=602,aet=102,crset=103,ad=24,cd=0,nasd=0)
> predict(mcfmodel, pred_dataglm, type="response")
      1
0.3189236
> pred_dataglm <- data.frame(mon=5,dm=2,uq="MQ",fn=3908,aet=129,crset=100,ad=160,cd=0,nasd=29)
> predict(mcfmodel, pred_dataglm, type="response")
      1
0.2283462
> pred_dataglm <- data.frame(mon=3,dm=28,uq="UA",fn=622,aet=95,crset=106,ad=42,cd=0,nasd=0)
> predict(mcfmodel, pred_dataglm, type="response")
      1
0.05702596
>
>
```